



# DATA SCIENCE

## PPGIA/PUCPR

Prof. Jean Paul Barddal



**REMINDERS**

## Reminders

- Today we will discuss dimensionality reduction
- Next week we will have our test!

# **DIMENSIONALITY REDUCTION**

## Dimensionality reduction

- In opposition to feature selection, dimensionality reduction techniques decrease the dimensionality of a problem by **combining** features
- There are different techniques to achieve this goal
- The most famous is Principal Component Analysis (PCA)

# **PRINCIPAL COMPONENT ANALYSIS**

## Variance and covariance

- Variance and covariance measure how “spread” a set of points are around their mean
- Variance is used for analyzing a single dimension
- Covariance measures how much each of the dimensions vary from the mean with respect to each other
- Covariance is measures between 2 dimensions to see if there is a relationship between them

## Covariance

- The covariance between 2 variables is computed by:

$$\text{Cov}(X, Y) = \frac{\sum_{i=1}^n (X_i - \bar{X})(Y_i - \bar{Y})}{(n - 1)}$$

- If you have more than 2 variables, you need to compute a covariance matrix



## Covariance

- The exact value is not as important as its sign
- A **positive** value indicates both variables increase or decrease together
- A **negative** value indicates that while one variable increases, the other decreases, or vice-versa
- If covariance is **zero**, the two variables are independent from each other

But what about correlation?

- Correlation allowed us to infer the same thing
- Why do we need covariance?
- Covariance is used to find relationships between variables in high-dimensional scenarios, where visualization is difficult

## Principal component analysis (PCA)

- PCA is a technique used to simplify a dataset
- It is a linear transformation that chooses a new coordinate system for the dataset such that:
  - **the greatest variance by any projection lies on the first axis:** the 1st principal component (eigenvector with the largest eigenvalue)
  - **the second greatest variance lies on the y axis (2nd PC), and so forth** (eigenvector with the second largest eigenvalue, etc)
- PCA can be used for reducing dimensionality by eliminating later principal components

## Steps to use PCA

- Normalize the data
  - Calculate the covariance matrix
  - Calculate the eigenvalues and eigenvectors
  - Choosing principal components
  - Forming a feature vector
  - Forming principal components
- 
- **All but the first of these steps are covered in scikit-learn's PCA implementation**

## PCA Limitations

- If the data does not follow a multidimensional normal (gaussian) distribution, the principal components extracted will be distorted

## Activity

- Let's run PCA on a dataset representing customers
- Each customer represents either a restaurant, a retail store, etc
- Let's analyze its principal components

**T-SNE**

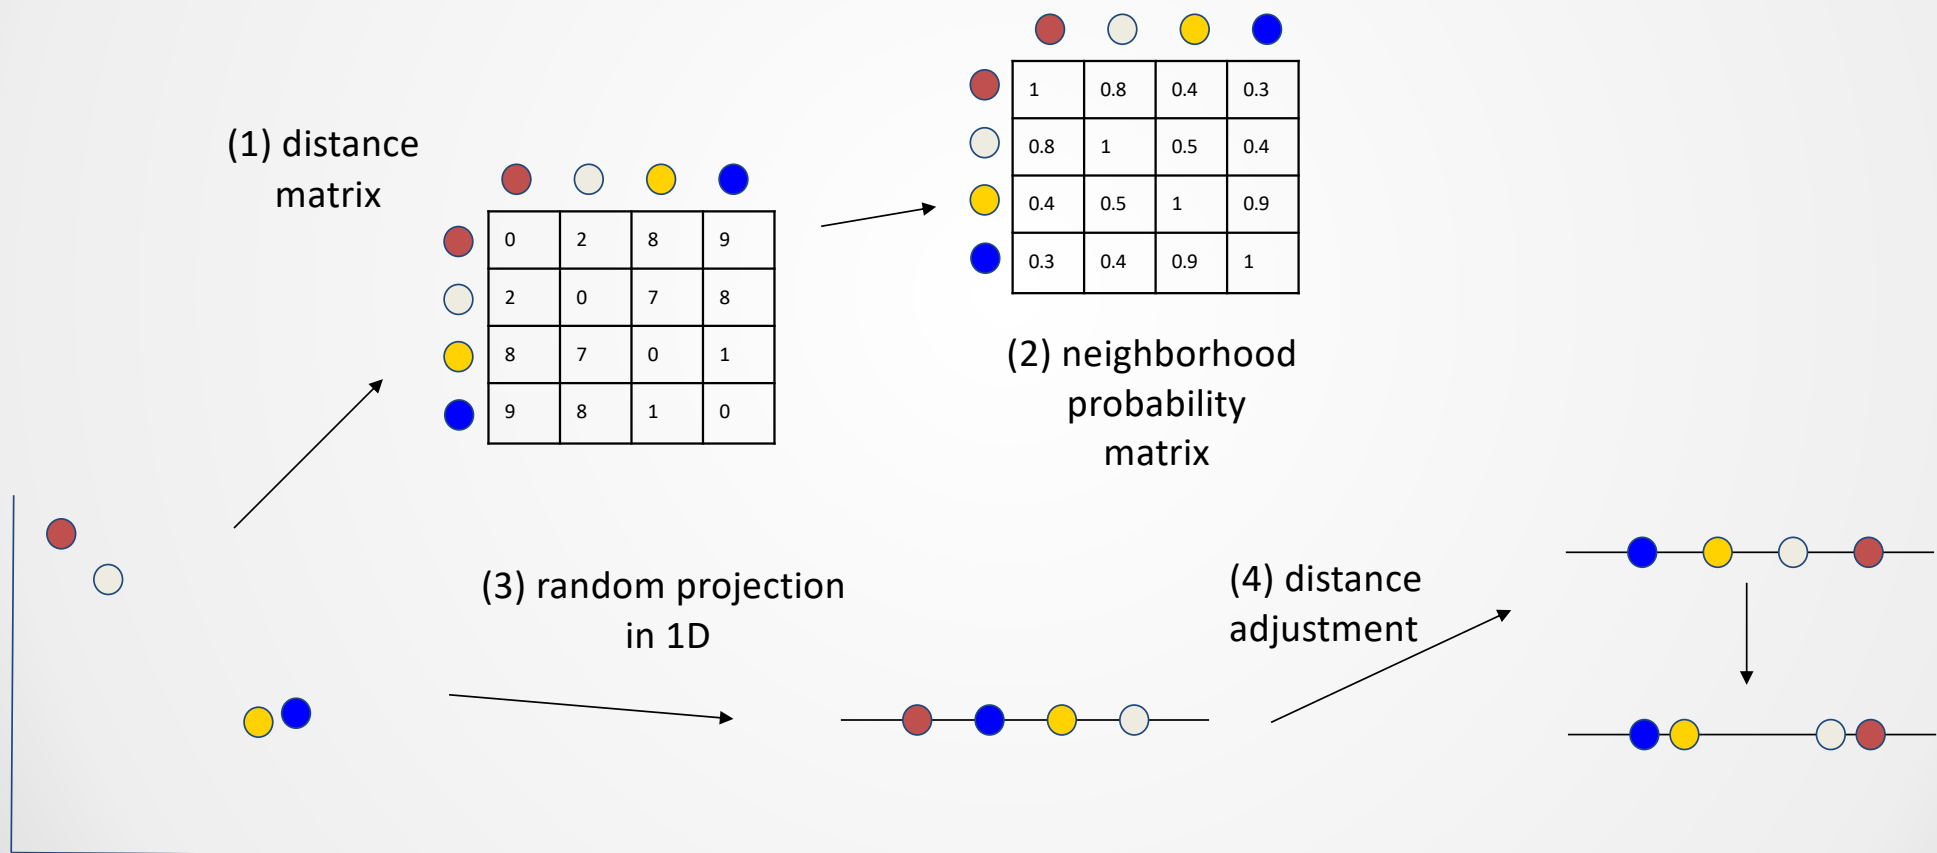
## t-Stochastic Neighbor Embedding

- Technique tailored for visualizing high-dimensional datasets
- How do we visualize data in 2D or 3D?
- Two goals:
  - Distance preservation
  - Neighbor preservation
- Unsupervised, but it helps uncovering interesting aspects of the data



# t-SNE overall idea

- Let's say we have a 2D problem we wish to visualize in 1D



# t-SNE

---

**Algorithm 1:** Simple version of t-Distributed Stochastic Neighbor Embedding.

---

**Data:** data set  $\mathcal{X} = \{x_1, x_2, \dots, x_n\}$ ,

cost function parameters: perplexity  $Perp$ ,

optimization parameters: number of iterations  $T$ , learning rate  $\eta$ , momentum  $\alpha(t)$ .

**Result:** low-dimensional data representation  $\mathcal{Y}^{(T)} = \{y_1, y_2, \dots, y_n\}$ .

**begin**

    compute pairwise affinities  $p_{j|i}$  with perplexity  $Perp$  (using Equation 1)

    set  $p_{ij} = \frac{p_{j|i} + p_{i|j}}{2n}$

    sample initial solution  $\mathcal{Y}^{(0)} = \{y_1, y_2, \dots, y_n\}$  from  $\mathcal{N}(0, 10^{-4}I)$

**for**  $t=1$  **to**  $T$  **do**

        compute low-dimensional affinities  $q_{ij}$  (using Equation 4)

        compute gradient  $\frac{\delta C}{\delta \mathcal{Y}}$  (using Equation 5)

        set  $\mathcal{Y}^{(t)} = \mathcal{Y}^{(t-1)} + \eta \frac{\delta C}{\delta \mathcal{Y}} + \alpha(t) (\mathcal{Y}^{(t-1)} - \mathcal{Y}^{(t-2)})$

**end**

**end**

# t-SNE

**Algorithm 1:** Simple version of t-Distributed Stochastic Neighbor Embedding.

**Data:** data set  $\mathcal{X} = \{x_1, x_2, \dots, x_n\}$ ,

cost function parameters: perplexity  $Perp$ ,

optimization parameters: number of iterations  $T$ , learning rate  $\eta$ , momentum  $\alpha(t)$ .

**Result:** low-dimensional data representation  $\mathcal{Y}^{(T)} = \{y_1, y_2, \dots, y_n\}$ .

**begin**

    compute pairwise affinities  $p_{j|i}$  with perplexity  $Perp$  (using Equation 1)

    set  $p_{ij} = \frac{p_{j|i} + p_{i|j}}{2n}$

    sample initial solution  $\mathcal{Y}^{(0)} = \{y_1, y_2, \dots, y_n\}$  from  $\mathcal{N}(0, 10^{-4}I)$

**for**  $t=1$  **to**  $T$  **do**

        compute low-dimensional affinities  $q_{ij}$  (using Equation 4)

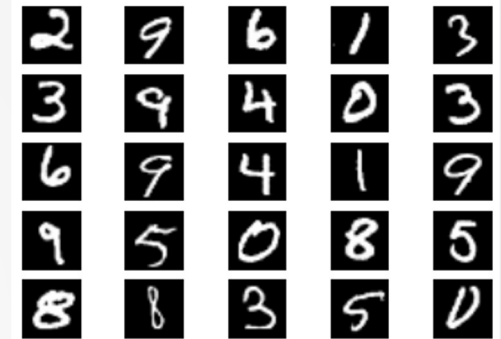
        compute gradient  $\frac{\delta C}{\delta \mathcal{Y}}$  (using Equation 5)

        set  $\mathcal{Y}^{(t)} = \mathcal{Y}^{(t-1)} + \eta \frac{\delta C}{\delta \mathcal{Y}} + \alpha(t) (\mathcal{Y}^{(t-1)} - \mathcal{Y}^{(t-2)})$

**end**

**end**

Random Sampling of MNIST



# t-SNE

**Algorithm 1:** Simple version of t-Distributed Stochastic Neighbor Embedding.

**Data:** data set  $\mathcal{X} = \{x_1, x_2, \dots, x_n\}$ ,

cost function parameters: perplexity  $Perp$ ,

optimization parameters: number of iterations  $T$ , learning rate  $\eta$ , momentum  $\alpha(t)$ .

**Result:** low-dimensional data representation  $\mathcal{Y}^{(T)} = \{y_1, y_2, \dots, y_n\}$ .

**begin**

    compute pairwise affinities  $p_{j|i}$  with perplexity  $Perp$  (using Equation 1)

    set  $p_{ij} = \frac{p_{j|i} + p_{i|j}}{2n}$

    sample initial solution  $\mathcal{Y}^{(0)} = \{y_1, y_2, \dots, y_n\}$  from  $\mathcal{N}(0, 10^{-4}I)$

**for**  $t=1$  **to**  $T$  **do**

        compute low-dimensional affinities  $q_{ij}$  (using Equation 4)

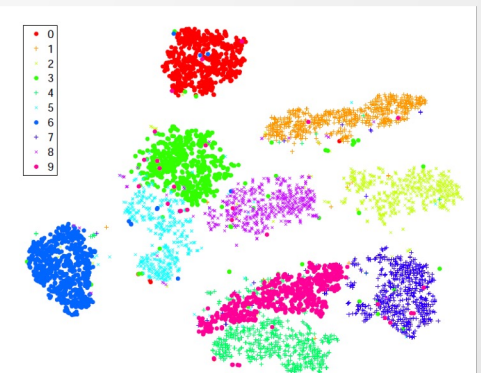
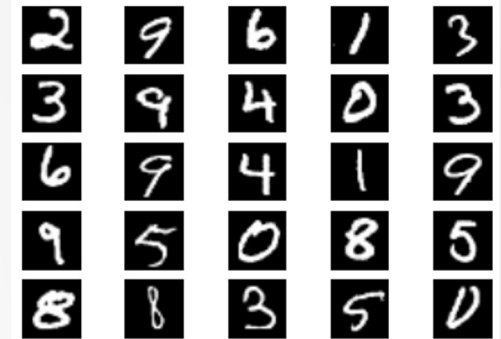
        compute gradient  $\frac{\delta C}{\delta \mathcal{Y}}$  (using Equation 5)

        set  $\mathcal{Y}^{(t)} = \mathcal{Y}^{(t-1)} + \eta \frac{\delta C}{\delta \mathcal{Y}} + \alpha(t) (\mathcal{Y}^{(t-1)} - \mathcal{Y}^{(t-2)})$

**end**

**end**

Random Sampling of MNIST



(a) Visualization by t-SNE.

# t-SNE

---

**Algorithm 1:** Simple version of t-Distributed Stochastic Neighbor Embedding.

---

**Data:** data set  $\mathcal{X} = \{x_1, x_2, \dots, x_n\}$ ,

cost function parameters: perplexity  $Perp$ ,

optimization parameters: number of iterations  $T$ , learning rate  $\eta$ , momentum  $\alpha(t)$ .

**Result:** low-dimensional data representation  $\mathcal{Y}^{(T)} = \{y_1, y_2, \dots, y_n\}$ .

**begin**

    compute pairwise affinities  $p_{j|i}$  with perplexity  $Perp$  (using Equation 1)

    set  $p_{ij} = \frac{p_{j|i} + p_{i|j}}{2n}$

    sample initial solution  $\mathcal{Y}^{(0)} = \{y_1, y_2, \dots, y_n\}$  from  $\mathcal{N}(0, 10^{-4}I)$

**for**  $t=1$  **to**  $T$  **do**

        compute low-dimensional affinities  $q_{ij}$  (using Equation 4)

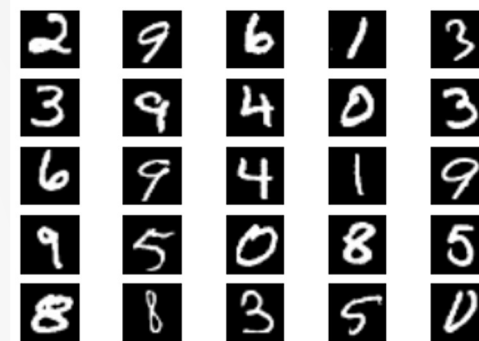
        compute gradient  $\frac{\delta C}{\delta \mathcal{Y}}$  (using Equation 5)

        set  $\mathcal{Y}^{(t)} = \mathcal{Y}^{(t-1)} + \eta \frac{\delta C}{\delta \mathcal{Y}} + \alpha(t) (\mathcal{Y}^{(t-1)} - \mathcal{Y}^{(t-2)})$

**end**

**end**

Random Sampling of MNIST



# t-SNE

## Algorithm 1: Simple version of t-Distributed Stochastic Neighbor Embedding.

**Data:** data set  $\mathcal{X} = \{x_1, x_2, \dots, x_n\}$ ,

cost function parameters: perplexity  $Perp$ ,

optimization parameters: number of iterations  $T$ , learning rate  $\eta$ , momentum  $\alpha(t)$ .

**Result:** low-dimensional data representation  $\mathcal{Y}^{(T)} = \{y_1, y_2, \dots, y_n\}$ .

**begin**

compute pairwise affinities  $p_{j|i}$  with perplexity  $Perp$  (using Equation 1)

set  $p_{ij} = \frac{p_{j|i} + p_{i|j}}{2n}$

sample initial solution  $\mathcal{Y}^{(0)} = \{y_1, y_2, \dots, y_n\}$  from  $\mathcal{N}(0, 10^{-4}I)$

**for**  $t=1$  **to**  $T$  **do**

compute low-dimensional affinities  $q_{ij}$  (using Equation 4)

compute gradient  $\frac{\delta C}{\delta \mathcal{Y}}$  (using Equation 5)

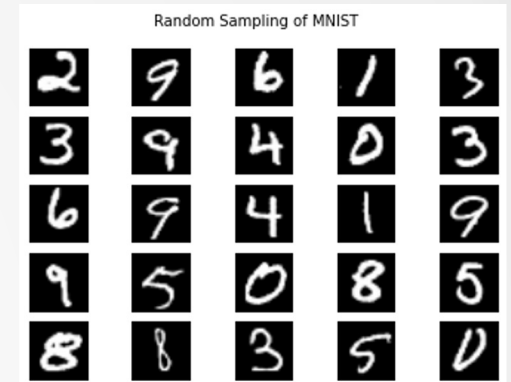
set  $\mathcal{Y}^{(t)} = \mathcal{Y}^{(t-1)} + \eta \frac{\delta C}{\delta \mathcal{Y}} + \alpha(t) (\mathcal{Y}^{(t-1)} - \mathcal{Y}^{(t-2)})$

**end**

**end**

**Perplexity** is the number of instances that we want to present the distances

Compute probabilities  $P$  that  $x_i$  and  $x_j$  are neighbors (based on Euclidian distance in **high-d** space)



# t-SNE

## Algorithm 1: Simple version of t-Distributed Stochastic Neighbor Embedding.

**Data:** data set  $\mathcal{X} = \{x_1, x_2, \dots, x_n\}$ ,

cost function parameters: perplexity  $Perp$ ,

optimization parameters: number of iterations  $T$ , learning rate  $\eta$ , momentum  $\alpha(t)$ .

**Result:** low-dimensional data representation  $\mathcal{Y}^{(T)} = \{y_1, y_2, \dots, y_n\}$ .

**begin**

    compute pairwise affinities  $p_{j|i}$  with perplexity  $Perp$  (using Equation 1)

    set  $p_{ij} = \frac{p_{j|i} + p_{i|j}}{2n}$

    sample initial solution  $\mathcal{Y}^{(0)} = \{y_1, y_2, \dots, y_n\}$  from  $\mathcal{N}(0, 10^{-4}I)$

**for**  $t=1$  **to**  $T$  **do**

        compute low-dimensional affinities  $q_{ij}$  (using Equation 4)

        compute gradient  $\frac{\delta C}{\delta y}$  (using Equation 5)

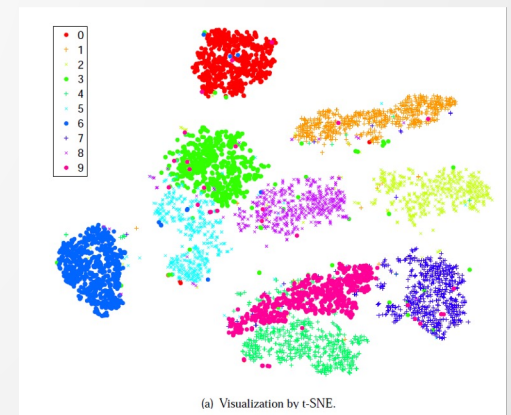
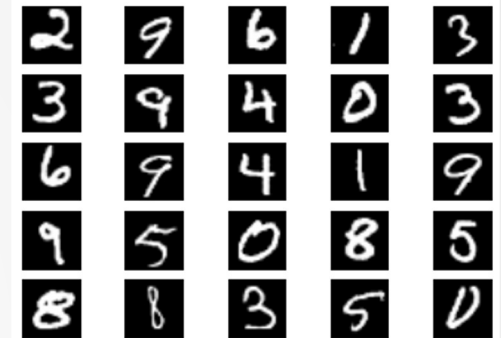
        set  $\mathcal{Y}^{(t)} = \mathcal{Y}^{(t-1)} + \eta \frac{\delta C}{\delta y} + \alpha(t) (\mathcal{Y}^{(t-1)} - \mathcal{Y}^{(t-2)})$

**end**

**end**

Compute probabilities  $Q$  that  $y_i$  and  $y_j$  are neighbors (corresponding to  $x_i, x_j$ )  
(based on Euclidian distance in **low-d** space)

Random Sampling of MNIST



# t-SNE

**Algorithm 1:** Simple version of t-Distributed Stochastic Neighbor Embedding.

**Data:** data set  $\mathcal{X} = \{x_1, x_2, \dots, x_n\}$ ,

cost function parameters: perplexity  $Perp$ ,

optimization parameters: number of iterations  $T$ , learning rate  $\eta$ , momentum  $\alpha(t)$ .

**Result:** low-dimensional data representation  $\mathcal{Y}^{(T)} = \{y_1, y_2, \dots, y_n\}$ .

**begin**

compute pairwise affinities  $p_{j|i}$  with perplexity  $Perp$  (using Equation 1)

set  $p_{ij} = \frac{p_{j|i} + p_{i|j}}{2n}$

sample initial solution  $\mathcal{Y}^{(0)} = \{y_1, y_2, \dots, y_n\}$  from  $\mathcal{N}(0, 10^{-4}I)$

**for**  $t=1$  **to**  $T$  **do**

compute low-dimensional affinities  $q_{ij}$  (using Equation 4)

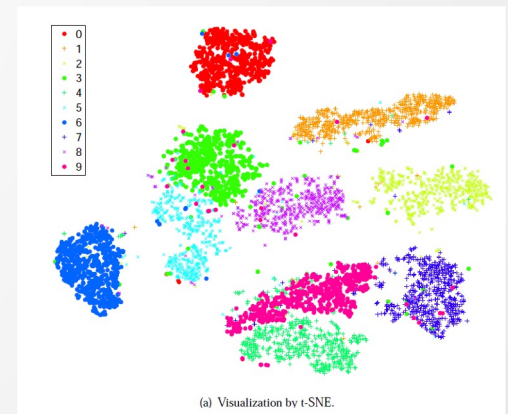
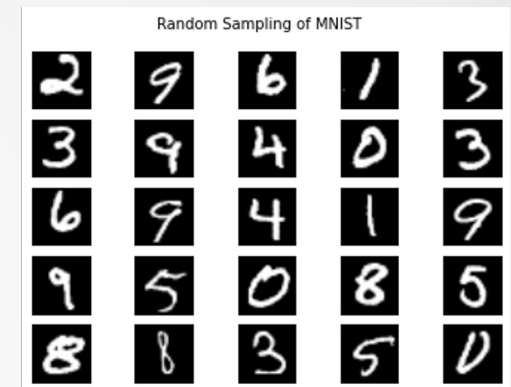
compute gradient  $\frac{\delta C}{\delta \mathcal{Y}}$  (using Equation 5)

set  $\mathcal{Y}^{(t)} = \mathcal{Y}^{(t-1)} + \eta \frac{\delta C}{\delta \mathcal{Y}} + \alpha(t) (\mathcal{Y}^{(t-1)} - \mathcal{Y}^{(t-2)})$

**end**

**end**

Key assumption is that the **high-d**  $P$  and the **low-d**  $Q$  probability distributions should be the same





# t-SNE

---

**Algorithm 1:** Simple version of t-Distributed Stochastic Neighbor Embedding.

---

**Data:** data set  $\mathcal{X} = \{x_1, x_2, \dots, x_n\}$ ,

cost function parameters: perplexity  $Perp$ ,

optimization parameters: number of iterations  $T$ , learning rate  $\eta$ , momentum  $\alpha(t)$ .

**Result:** low-dimensional data representation  $\mathcal{Y}^{(T)} = \{y_1, y_2, \dots, y_n\}$ .

**begin**

    compute pairwise affinities  $p_{j|i}$  with perplexity  $Perp$  (using Equation 1)

    set  $p_{ij} = \frac{p_{j|i} + p_{i|j}}{2n}$

    sample initial solution  $\mathcal{Y}^{(0)} = \{y_1, y_2, \dots, y_n\}$  from  $\mathcal{N}(0, 10^{-4}I)$

**for**  $t=1$  **to**  $T$  **do**

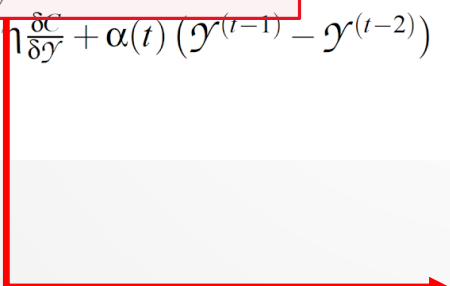
        compute low-dimensional affinities  $q_{ij}$  (using Equation 4)

        compute gradient  $\frac{\delta C}{\delta \mathcal{Y}}$  (using Equation 5)

        set  $\mathcal{Y}^{(t)} = \mathcal{Y}^{(t-1)} + \eta \frac{\delta C}{\delta \mathcal{Y}} + \alpha(t) (\mathcal{Y}^{(t-1)} - \mathcal{Y}^{(t-2)})$

**end**

**end**



Find a **low-d** map that minimizes the difference between the **P** (high-d) and **Q** (low-d) distributions

(if  $x_i, x_j$  has high probability of being neighbors in **high-d**, then  $y_i, y_j$  should have high probability in **low-d**)

# t-SNE

**Algorithm 1:** Simple version of t-Distributed Stochastic Neighbor Embedding.

**Data:** data set  $\mathcal{X} = \{x_1, x_2, \dots, x_n\}$ ,

cost function parameters: perplexity  $Perp$ ,

optimization parameters: number of iterations  $T$ , learning rate  $\eta$ , momentum  $\alpha(t)$ .

**Result:** low-dimensional data representation  $\mathcal{Y}^{(T)} = \{y_1, y_2, \dots, y_n\}$ .

**begin**

    compute pairwise affinities  $p_{j|i}$  with perplexity  $Perp$  (using Equation 1)

    set  $p_{ij} = \frac{p_{j|i} + p_{i|j}}{2n}$

    sample initial solution  $\mathcal{Y}^{(0)} = \{y_1, y_2, \dots, y_n\}$  from  $\mathcal{N}(0, 10^{-4}I)$

**for**  $t=1$  **to**  $T$  **do**


        compute low-dimensional affinities  $q_{ij}$  (using Equation 4)

        compute gradient  $\frac{\delta C}{\delta \mathcal{Y}}$  (using Equation 5)

        set  $\mathcal{Y}^{(t)} = \mathcal{Y}^{(t-1)} + \eta \frac{\delta C}{\delta \mathcal{Y}} + \alpha(t) (\mathcal{Y}^{(t-1)} - \mathcal{Y}^{(t-2)})$

**end**

**end**



We will minimize the difference between the **high-d** and **low-d** maps using **gradient descent**

## t-SNE details

- Details on t-SNE can be found at the original paper
- <https://jmlr.org/papers/volume9/vandermaten08a/vandermaten08a.pdf>

Journal of Machine Learning Research 9 (2008) 2579-2605

Submitted 5/08; Revised 9/08; Published 11/08

### Visualizing Data using t-SNE

**Laurens van der Maaten**

*TiCC*

*Tilburg University*

*P.O. Box 90153, 5000 LE Tilburg, The Netherlands*

LVDMAATEN@GMAIL.COM

**Geoffrey Hinton**

*Department of Computer Science*

*University of Toronto*

*6 King's College Road, M5S 3G4 Toronto, ON, Canada*

HINTON@CS.TORONTO.EDU

**Editor:** Yoshua Bengio

# t-SNE

- In opposition to PCA, t-SNE is **not** parametric
- This means that we cannot learn a manifold representation from a dataset and apply it to another dataset
  - Therefore, this cannot be used as a dimensionality reduction technique between training and test data
- There is, however, a parametric version available at:

---

## Learning a Parametric Embedding by Preserving Local Structure

---

Laurens van der Maaten  
TiCC, Tilburg University  
P.O. Box 90153, 5000 LE Tilburg, The Netherlands  
lvdmaaten@gmail.com